

---

---

# CS205 Project Proposal

## Team 06

---

---

Zhecheng Yao, Yixian Gan, Rebecca Qiu, Lucy Li

---

---

# FaceX: Face Contour Recognition via Shape Regression

## Face recognition

- Robustness: variations in face pose, expression, and lighting conditions
- Flexibility: detect a wide range of facial landmarks (e.g. eye, eyebrow, lips)
- Efficiency: relative lean and simple model to perform real time analysis

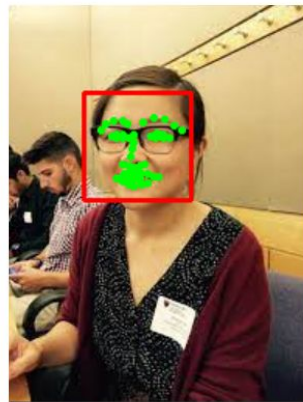
*Our facial regression model aims to predict the position of facial landmarks, using a boosted regressor to make progressive inference, minimizing alignment error at each step.*

## Challenges:

- current version w/o parallelism has significant lagging in performance
- 1.5 seconds between face update and face recognition
- ~1 hour time requirement to train on a small dataset of 60  $250 \times 250$  pixel image

# Two Parts

- Face recognition w/ trained model
- Training model




# Model and Data


## Basic Pipeline of Recognition (FaceX):

1. Load an image frame from camera, convert it to grayscale, maps onto a unit square and detect faces present.
2. For each detected face, draw a red rectangle to bound it.
3. align the initial facial landmarks to the face rectangle

### Potential Optimizations



3. With OpenMP, split the input landmark into multiple batches and process each subset in parallel (reduction)



# Model and Data


## Basic Pipeline of Recognition (FaceX):

1. Load an image frame from camera, convert it to grayscale, maps onto a unit square and detect faces present.
2. For each detected face, draw a red rectangle to bound it.
3. Align the initial facial landmarks to the face rectangle
4. Apply trained regression model using pixel within the rectangle and initial facial landmarks as inputs to predict refined landmark locations.
5. Compute prediction medians as final landmarks' coordinates.
6. Draw green dots the final landmarks and display the aligned image.


Data source:

Images and videos taken on our phone or camera


## Potential Optimizations



3. With OpenMP, split the input landmark into multiple batches and process each subset in parallel (reduction)



4. Use OpenMP to distribute the workload across multiple threads in the prediction stage



# Training Model

## FaceX-Train:

1. Creates a set of initial shapes for testing.
2. Maps the landmarks onto the unit square.
3. Augments Training data
4. Perform **gradient boosting regression** - two levels cascade, each iteration involves:

Randomly sampling  $P$  pixels, computing  $P^2$  "**pixel-difference features**"

Features are indexed based on how far they are from local landmark

★ Feature selection based on **correlations** - select  $F$  of the  $P^2$  features


**Normalize regression target**, aligning the target shape to the mean shape across all training dataset

★ Perform model compression using **orthogonal matching pursuit (OMP)**:

5. Returns the training parameters, training data, and test initial shapes.

★ Indicates Bottleneck in Sequential Version of the Code

{ Every step enclosed in brackets is individually parallelizable }



## Potential Optimizations

- ★ In computing covariance, use vectorization (AVX 2 instructions), or OpenMP to parallelize the implementation
- ★ Using MPI, distribute the input shape across multiple processes

# Profiling and Bottleneck

[illegible]

# Profiling and Bottleneck

```
double Covariance(double *x, double *y, const int size)
{
    double a = 0, b = 0, c = 0;
    for (int i = 0; i < size; ++i)
    {
        a += x[i];
        b += y[i];
        c += x[i] * y[i];
    }

    return c / size - (a / size) * (b / size);
}
```

Current performance:

- size=13466 (# of training sample)
- 4\*size per function call
- 400 function call per regressor
- 1000 regressor in total

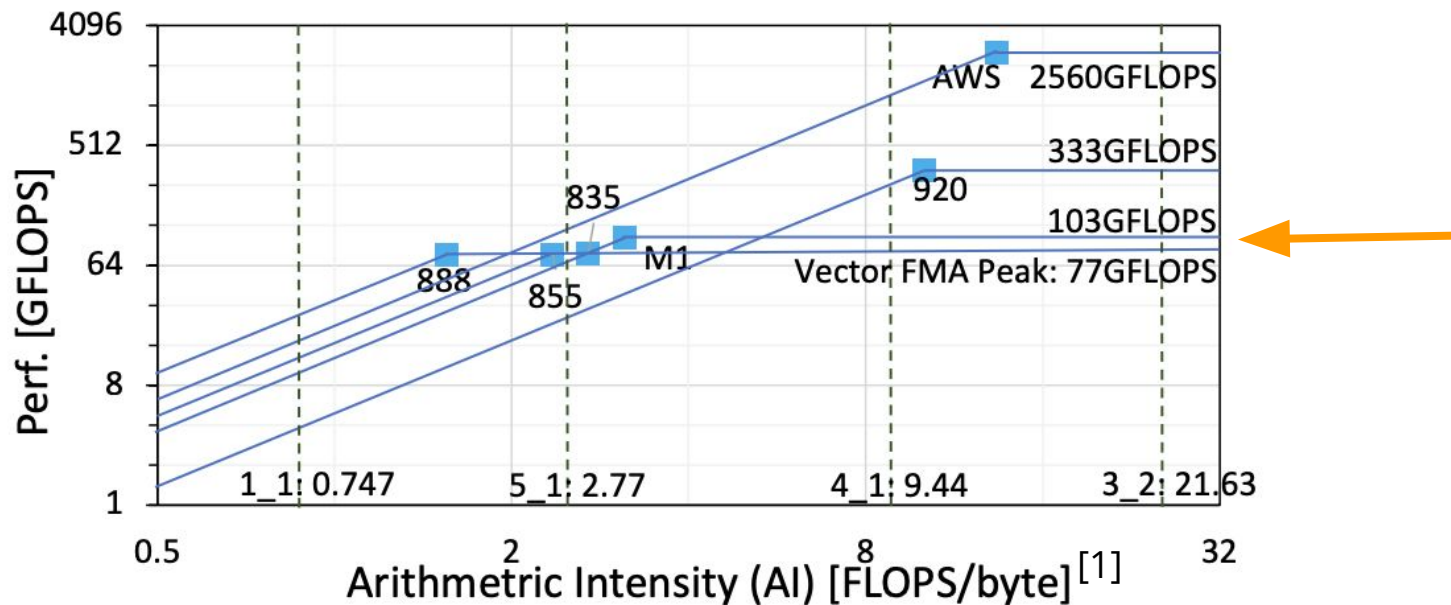
= 21.54 GFlop

Time  $\approx$  1000s

**0.02 GFlop/s**



# Roofline Analysis



[1]<https://www.osti.gov/pages/servlets/purl/1863284>

# Parallelism Offered by the Project

## Shared Memory Parallelism (OpenMP)

There are lots of For-loops to be parallelize

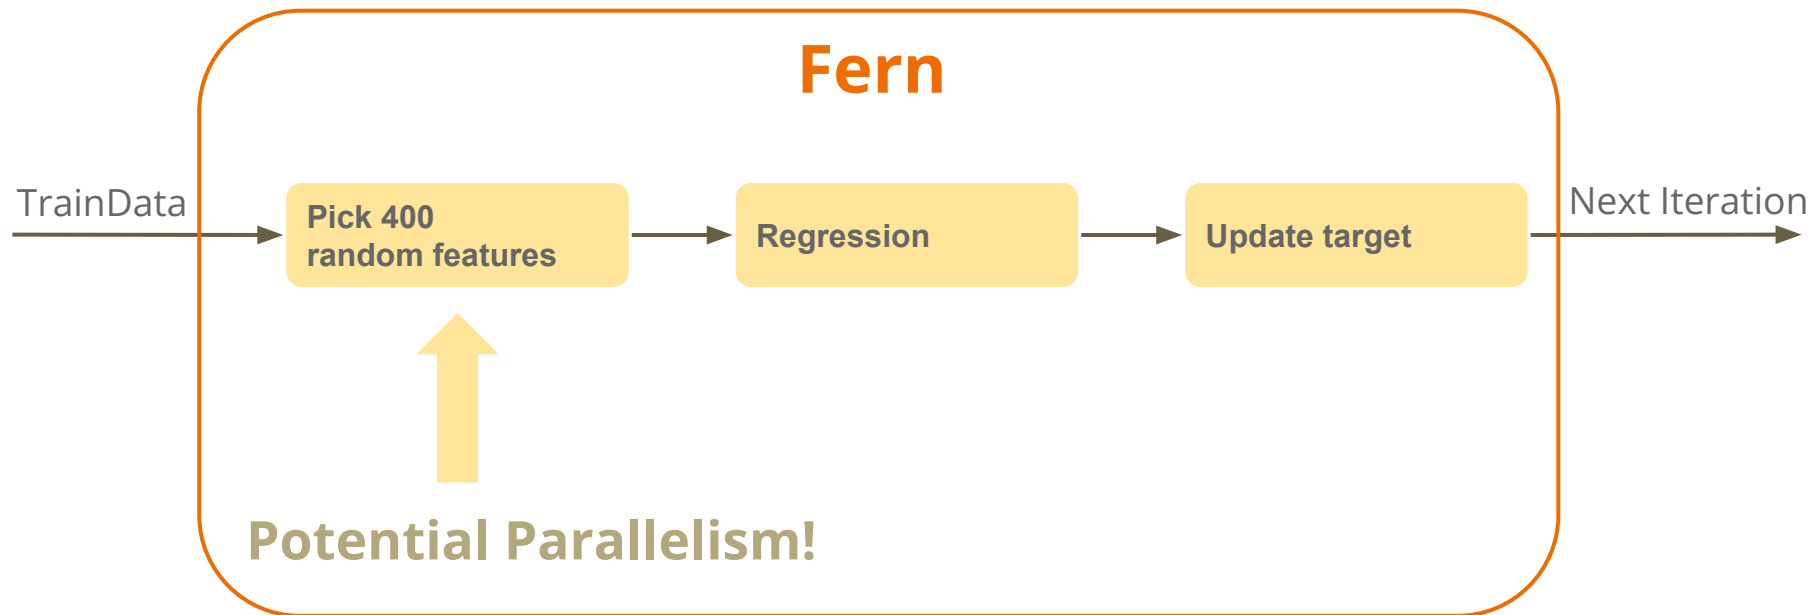
- `utils_train.cpp::Covariance()`
- `regressor_train.cpp::Regress()`
- `regressor_train.cpp::CompressFerns()`
- ...

```
for (int i = 0; i < pixels_val.cols; ++i)
{
    Transform t = Procrustes(training_data[i].init_shape, mean_shape);
    vector<cv::Point2d> offsets(training_parameters_.P);
    for (int j = 0; j < training_parameters_.P; ++j)
        offsets[j] = pixels_[j].second;
    t.Apply(&offsets, false);

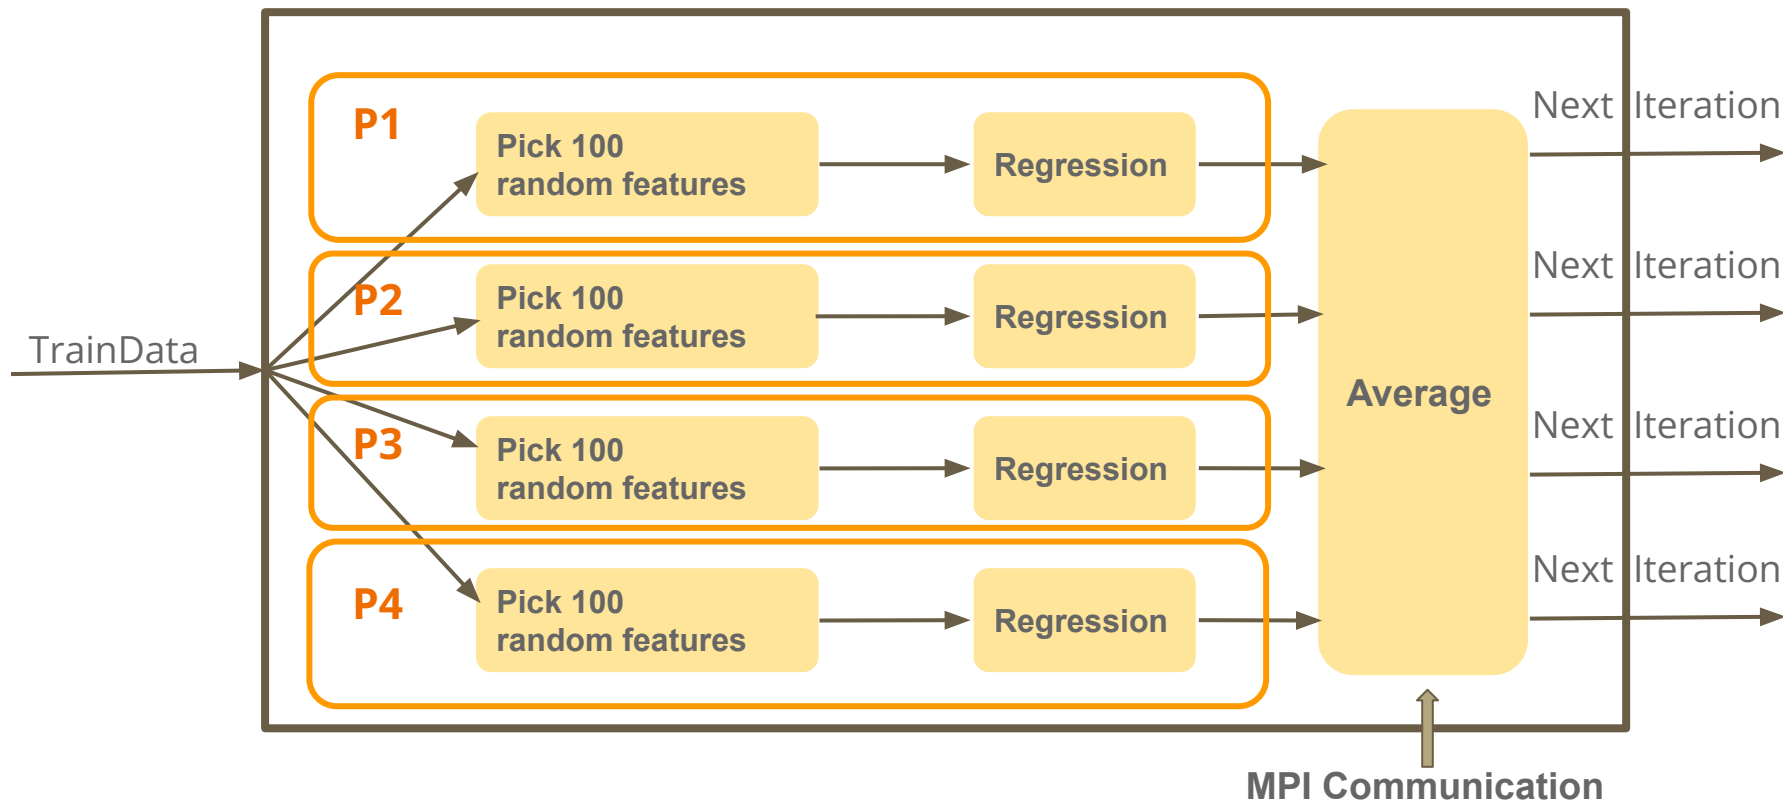
    for (int j = 0; j < training_parameters_.P; ++j)
    {
        cv::Point pixel_pos = training_data[i].init_shape[pixels_[j].first]
            + offsets[j];
        if (pixel_pos.inside(cv::Rect(0, 0,
            training_data[i].image.cols, training_data[i].image.rows)))
        {
            pixels_val.at<double>(j, i) =
                training_data[i].image.at<uchar>(pixel_pos);
        }
        else
            pixels_val.at<double>(j, i) = 0;
    }
}
```

# Parallelism Offered by the Project

## Distributed Memory Parallelism (MPI)



# Parallelism Offered by the Project



# Team 06

- Zhecheng Yao: [zhechengyao@g.harvard.edu](mailto:zhechengyao@g.harvard.edu)
- Yixian Gan: [ygan@g.harvard.edu](mailto:ygan@g.harvard.edu)
- Rebecca Qiu: [zqiu1@fas.harvard.edu](mailto:zqiu1@fas.harvard.edu)
- Lucy Li: [cli@fas.harvard.edu](mailto:cli@fas.harvard.edu)