


Parallelizing Gradient Boosted Regression for Facial Landmark Recognition

Zhecheng Yao, Yixian Gan, Rebecca Qiu, Lucy Li
May 2023




Harvard John A. Paulson
School of Engineering
and Applied Sciences

Contents

Introduction & Research Problem

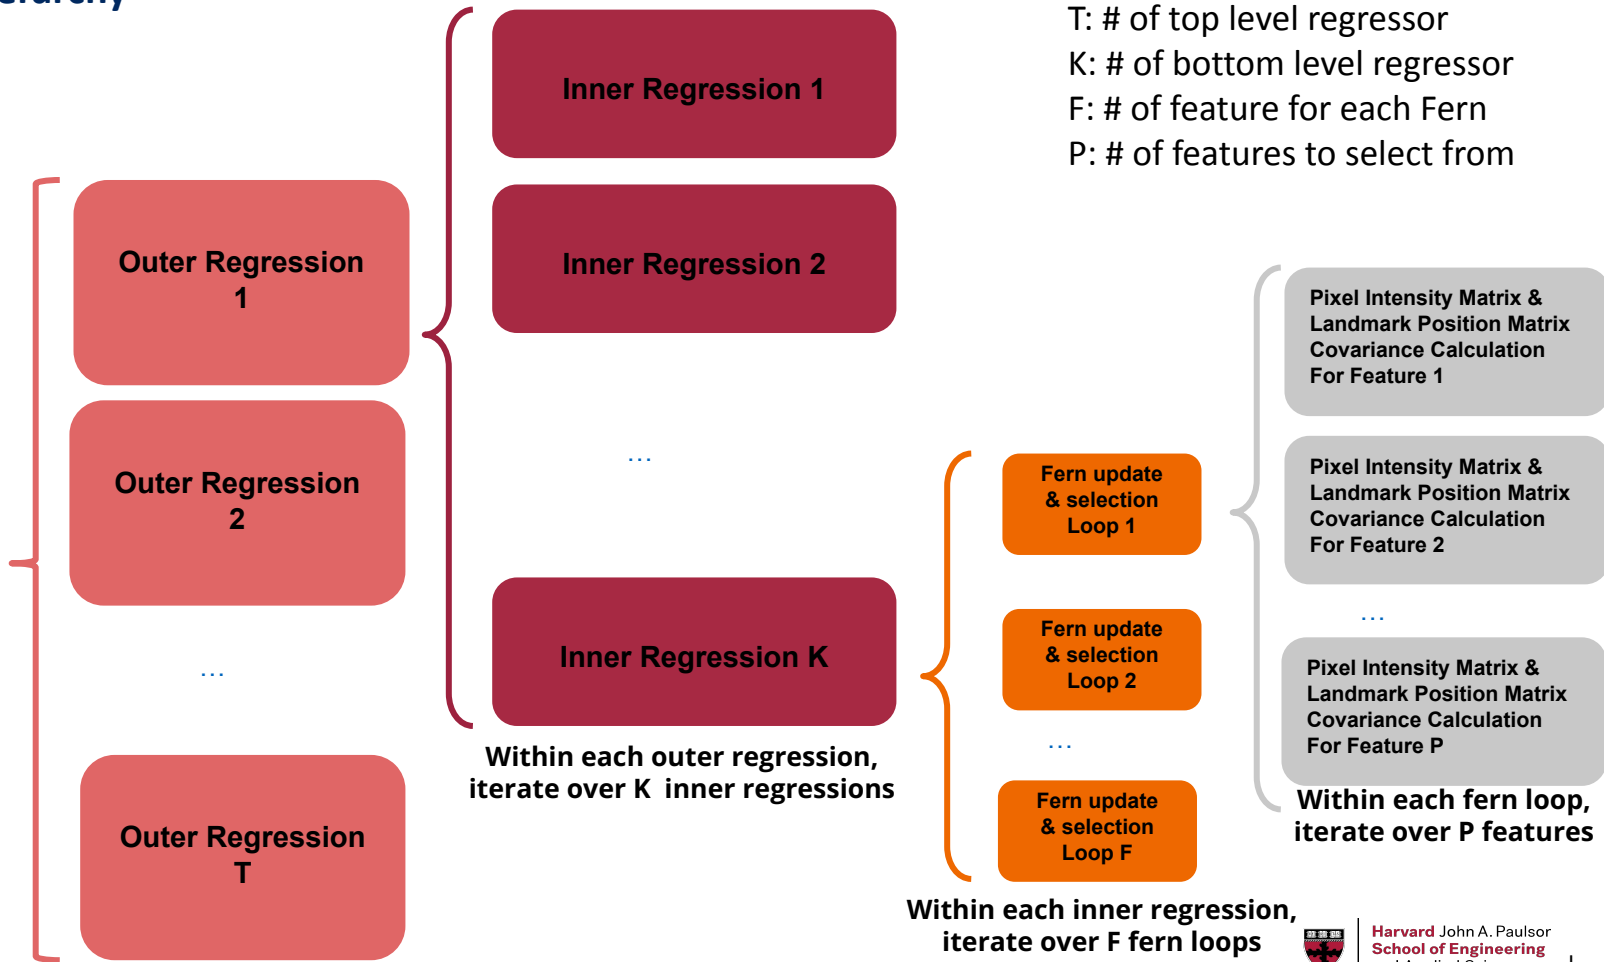
Methods & Algorithms

Results & Validation



Regression Hierarchy

T Iterations



T: # of top level regressor
 K: # of bottom level regressor
 F: # of feature for each Fern
 P: # of features to select from

Within each inner regression,
 iterate over F fern loops

Sequential Baseline



Baseline

Training Time ~ 200s

Data Size: 4000 images
Image Dimension: 250x250
CPU: Intel Xeon E5-2683 v4
Number of Regressors: 1

Objective

Increase training efficiency by implementing parallelism.




Contents

Introduction & Research Problem


Parallel Design & Algorithms

Results & Validation




Get Training Data Sequential


Work Logic



colored image of any size



grayscale matrix




four coordinates of face rectangle

vector $\{ (0,1), (0,4), (4,1), (4,4) \}$ 

stored under a "DataPoint" struct

Then all DataPoints are stored in a vector of DataPoints




facial landmark coordinates

(same landmark count for all input images)

vector $\{ (1.5,1.5), (2.5,3.5), (3.5,1.5) \}$ 

Get Training Data MPI



Rank 0 Image read

DataPoint 0

grayscale matrix 0


face rectangle vector 0

landmark vector 0

...

Design choice:
regressor parallelism
on feature-level
not data-level

serial code has m regressor
each regressor need all S images' data
for MPI code with rank count n
we still want m regressors
each train with S images
-> need communication to share image data



Rank 1 Image read

DataPoint 1

grayscale matrix 1

face rectangle vector 1



landmark vector 1

...

Rank 2 ...

...

communication
to distribute partial read data



Design Breakdown

Step 1 Each Rank Load Assigned Images

Read info text file

Text lines w/ data paths assigned to each rank

Each rank do parallel read in of images, face rectangle & facial landmarks

Step 2 Gather All Images to Rank 0

Serialize local DataPoints into array for send buffer

Gather sizes of each rank's send buffer to rank 0 and calculate displacements

Gatherv collect all image data

Deserialize rank 0 receive buffer into complete DataPoints

Step 3 Rank 0 Scatter Images to all ranks

rank 0 serialize complete DataPoints into array for send buffer

Broadcast send buffer size to all ranks so each rank can prepare properly sized receive buffer

Broadcast complete DataPoints

Each rank deserialize complete DataPoints and update DataPoints

DataPoint 1

grayscale matrix

```
[150 100 100 100 150
150 50 255 50 150
150 255 255 255 150
150 255 50 255 150
150 150 150 150 150]
```

face rectangle vector

{ (0,1), (0,4), (4,1), (4,4) }

landmark vector

{ (1.5,1.5), (2.5,3.5), (3.5,1.5) }

Serialize



Deserialize



Send/Recv Buffer

```
{ 5, 5, // matrix size
150 , 100, 100.... // first row in matrix
150, 50 .... // second and other rows
4, 2 // face rectangle vector size, each has x&y
0, 1, 0, 4 .... // face rectangle vector values
3, 2, 1.5, 1.5 .... // similar logic for landmark vector ...}
```

DataPoint 2, 3...



Parallel Result

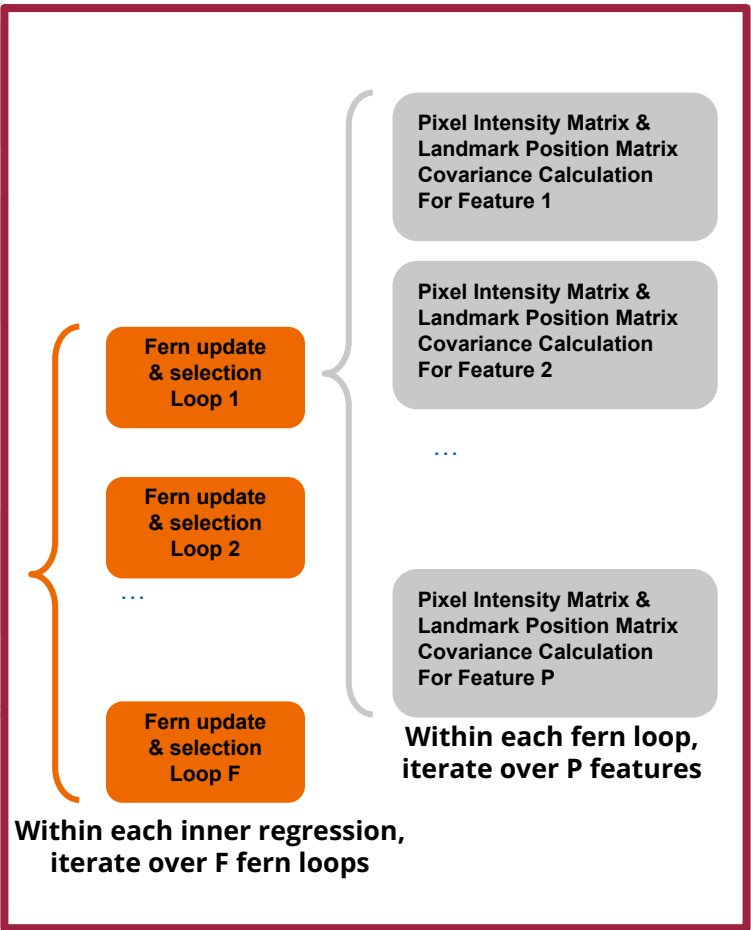
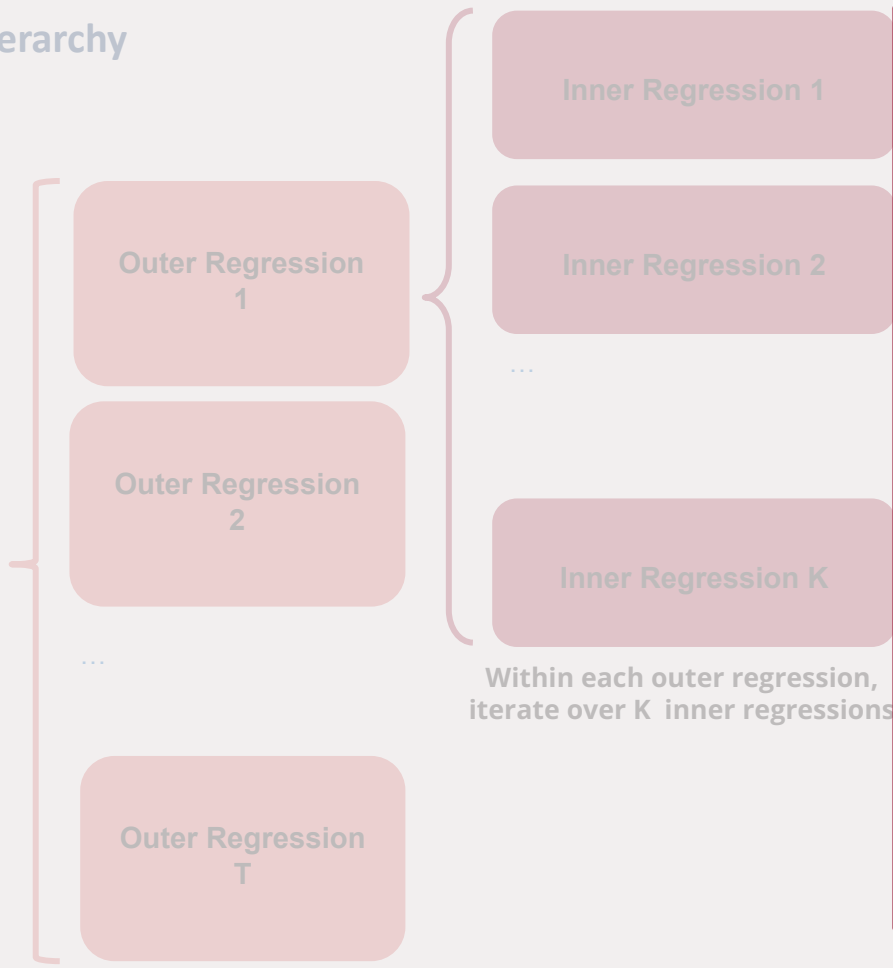


● Highlights

- More ranks provide better Speedup (~2-3x for 1-4 ranks)
- More Images take longer time to process
- Serialization/deserialization takes toll on performance

Regression Hierarchy

T Iterations




Fern Regress Sequential

T_{f0} = Time to select feature 0

T_{i0} = Time to compute feature of image 0

T_{o0} = Time to compute outputs 0

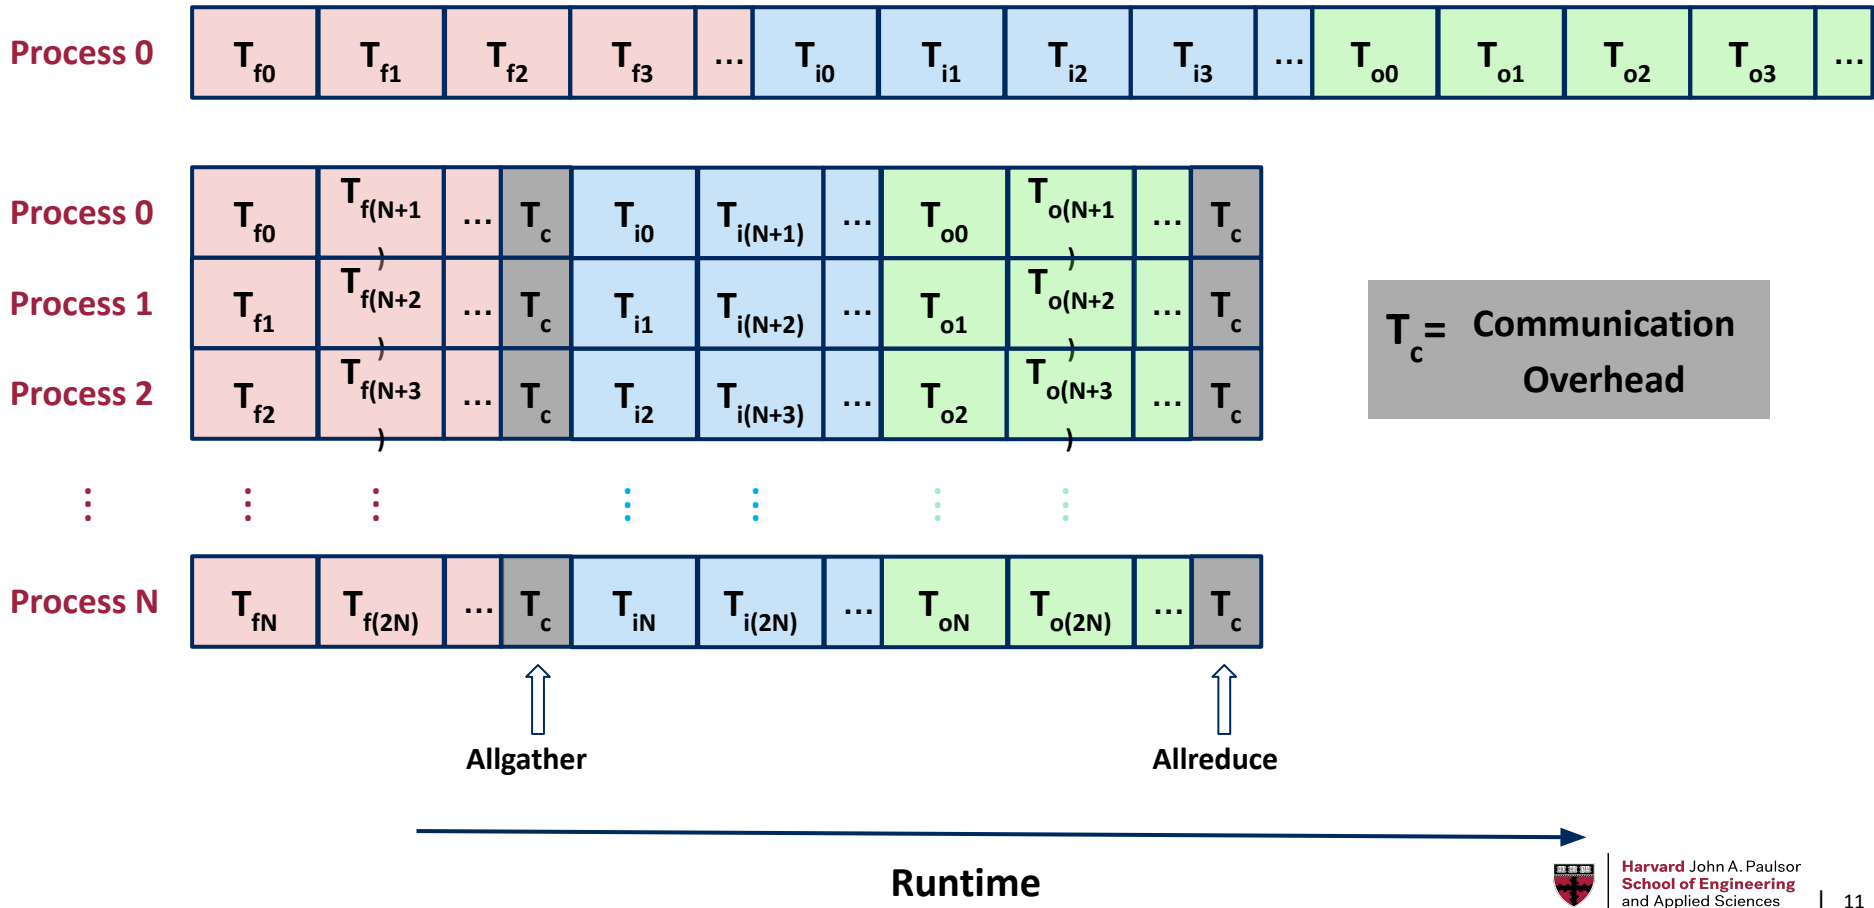
Process 0





Runtime



Fern Regress MPI



Parallel Result




- Highlights
 - More ranks give better performance up to a threshold
 - Larger improvements for larger problem size (# of training images)
 - Efficiency depends on training hyper-parameter

Orthogonal Matching Pursuit - OpenCV vs Eigen

Target: Reconstructing of noisy result vector from the Fern regressions $\hat{y} \in \mathbb{R}^{N \times 1}$.

1. Samples a matrix B



$$B \in \mathbb{R}^{N \times M}$$

2. $residual \leftarrow \hat{y}$


3. OMP chooses the best Q candidate basis vector in matrix B maximizing dot product (Greedy)

$$B_j \leftarrow \arg \max_{B_{j'} \in B} |\langle residual, B_{j'} \rangle|$$



residual

\times



B_j

- **FLOPS:**
(N multiplications + N - 1 additions) M (columns)
Q (iterations)

**Best Basis Vector
Selection**



Orthogonal Matching Pursuit - OpenCV vs Eigen



3

$$B_j \leftarrow \arg \max_{B_{j'} \in B} |\langle residual, B_{j'} \rangle|$$

\\ Native OpenCV implementation

```
double current_value = abs(static_cast<cv::Mat>(residual.t() * base.col(j)).at<double>(0));
```

\\ Eigen Implementation


```
Eigen::MatrixXd column = baseEigen.col(j);
```

```
double current_value = (residual.transpose() * column).cwiseAbs().sum();
```



residual

×



B_j

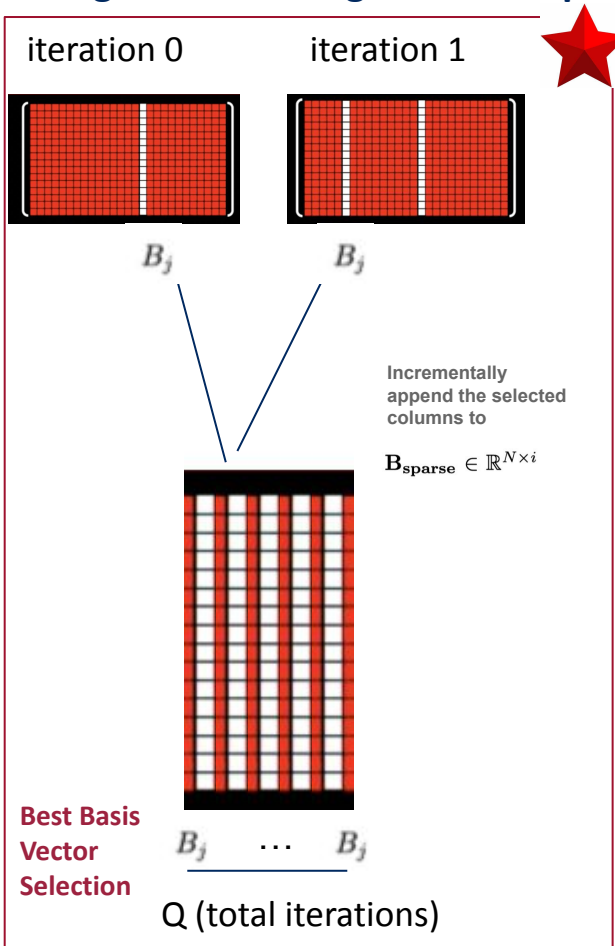
- **FLOPS:**
(N multiplications + N - 1 additions) M
(columns) Q (iterations)

**Best Basis Vector
Selection**



Orthogonal Matching Pursuit - OpenCV vs Eigen

3



4

Dimensions

$$B_{\text{sparse}}^T \in \mathbb{R}^{i \times N}$$

$$B_{\text{sparse}} \in \mathbb{R}^{N \times i}$$

$$\hat{y} \in \mathbb{R}^{N \times 1}$$

Optimization Problem

$$x_i \leftarrow \arg \min_{x_i} \|\hat{y} - B_{\text{sparse}} x_i\|$$

$$x_i = (B_{\text{sparse}}^T B_{\text{sparse}})^{-1} B_{\text{sparse}}^T \hat{y}.$$

FLOPS in the i-th iteration

FLOPS in the Q total iterations

$B_{\text{sparse}}^T B_{\text{sparse}}$	$(2N - 1)i^2$	$(2N - 1)(Q/6)(Q + 1)(2 * Q + 1)$
$B_{\text{sparse}}^T \hat{y}$	$(2 * N - 1) * i$	$(Q)(Q + 1)/2 * (2 * N - 1)$
SVD	$13 * i^3$	$13 * 0.25 * (Q^4 + 2 * Q^3 + Q^2)$

Solving the Linear Systems

5

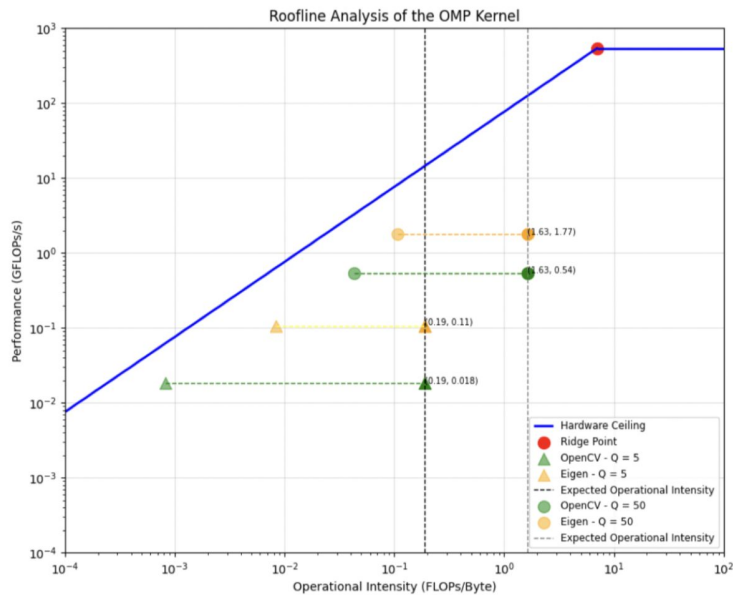
$$residual \leftarrow residual - B_{\text{sparse}} x_i$$

FLOPS in the Q total iterations



$$((Q)(Q + 1)/2 - 1)N + N * Q$$



Orthogonal Matching Pursuit - OpenCV vs Eigen



(a) Roofline Analysis



(b) Scaling Analysis

Figure 4: Performance Analysis of the OMP Kernel


Note: **OpenCV** in green, **Eigen** in yellow

Faint triangles and circles are PAPI measured
theoretical OI is darker circles and triangles


- Total L1 cache misses for $Q = 5$ with OpenCV is **26,327**
- Total L1 cache misses for $Q = 5$ with the Eigen library is **7554**

Covariance Sequential Version

2 arrays both of size n



Within i-th step



At the very end, return

$$\frac{\text{sum_xy}}{n} - \frac{\text{sum_x}}{n} \times \frac{\text{sum_y}}{n}$$



Covariance ISPC Version

2 arrays both of size n



use gang size of 4

Time/ n Iterations

For each gang j at step i

$$\text{sum_x} += (i-1) \times 4 + j \text{ th}$$

$$\text{sum_y} += (i-1) \times 4 + j \text{ th}$$



$$\text{sum_xy} += (i-1) \times 4 + j \text{ th} \times (i-1) \times 4 + j \text{ th}$$

At the very end, return

$$\frac{\text{reduced sum_xy}}{n} - \frac{\text{reduced sum_x}}{n} \times \frac{\text{reduced sum_y}}{n}$$



Covariance - ISPC Performance Analysis



- Highlights
 - ISPC at same level of magnitude as nominal peak performance
 - Could push further with loop unrolling to saturate all registers
 - ~4x time reduction for a medium sized data input (O2 flag)



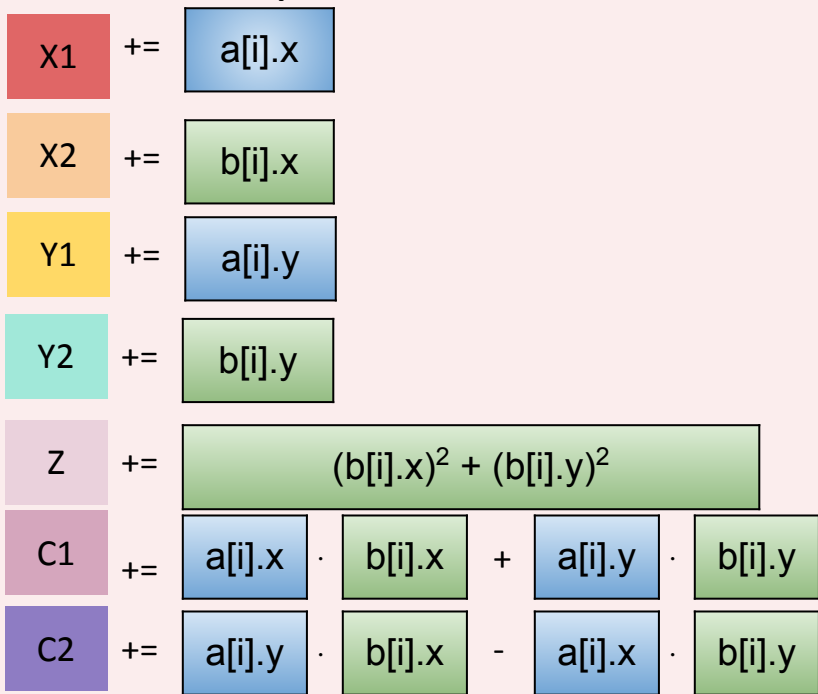
Procrustes: find the optimal linear transformation between two sets of DataPoints

2 arrays of size N

Point2d Object a	$a[0] = (a[0].x, a[0].y)$	$a[1] = (a[1].x, a[1].y)$...	$a[N] = (a[N].x, a[N].y)$
Point2d Object b	$b[0] = (b[0].x, b[0].y)$	$b[1] = (b[1].x, b[1].y)$...	$b[N] = (b[N].x, b[N].y)$

 Time

Within the for loop:



At the very end:

Matrix A =

$$\begin{bmatrix} X2 & -Y2 & N & 0 \\ Y2 & X2 & 0 & N \\ Z & 0 & X2 & Y2 \\ 0 & Z & -Y2 & X2 \end{bmatrix}$$

Matrix b =

$$\begin{bmatrix} X1 & Y1 & C1 & C2 \end{bmatrix}$$

Return:

Transformation Vector = $A.inv() \cdot b$



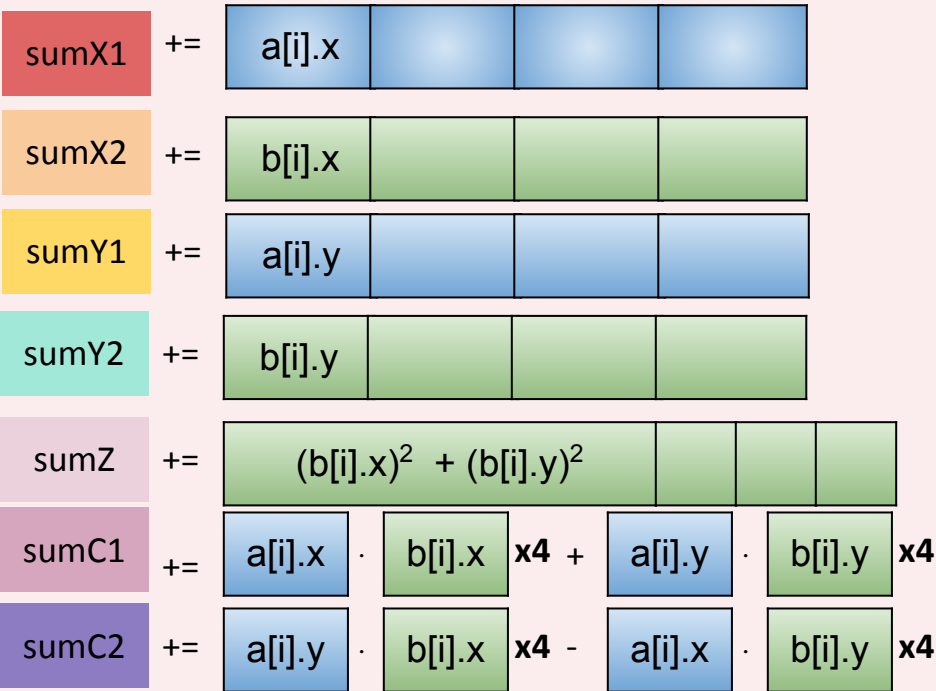
Procrustes - AVX2

2 arrays of size N

Point2d Object a	$a[0] = (a[0].x, a[0].y)$	$a[1] = (a[1].x, a[1].y)$...	$a[N] = (a[N].x, a[N].y)$
Point2d Object b	$b[0] = (b[0].x, b[0].y)$	$b[1] = (b[1].x, b[1].y)$...	$b[N] = (b[N].x, b[N].y)$

→ Time

Within the for loop : unroll by 2



At the very end:

Matrix A =

$$\begin{bmatrix} X2 & -Y2 & N & 0 \\ Y2 & X2 & 0 & N \\ Z & 0 & X2 & Y2 \\ 0 & Z & -Y2 & X2 \end{bmatrix}$$

Matrix b =



$$\begin{bmatrix} X1 & Y1 & C1 & C2 \end{bmatrix}$$

Return:


Transformation Vector = $A.inv() \cdot b$



Procrustes - AVX2



**with -O3 flag
2x Speedup**




**without -O3 flag
4x Speedup**

Contents




Introduction & Research Problem

Parallel Design & Algorithms

Results & Validation



Performance Benchmarks and Strong Scaling



Optimal Number of Parallel Processes: 5

Sequential vs. Parallelized

Process Count	1	2	3	4	5	6	8	10	12	14	16	baseline	speedup
250 Images	14.70	13.56	13.02	12.89	11.71	11.74	12.19	12.10	12.22	12.31	12.21	51.54	4.41
500 Images	21.31	19.01	18.01	17.67	16.32	17.09	17.43	17.74	17.77	17.60	17.54	59.85	3.67
1000 Images	44.51	34.68	31.32	31.10	27.40	27.66	28.06	28.68	29.18	29.41	29.44	81.93	2.99
2000 Images	84.26	66.70	58.75	57.14	49.52	51.93	50.40	51.93	50.99	52.23	52.57	120.86	2.44
4000 images	158.72	125.12	113.11	106.43	91.52	93.68	95.86	94.76	94.15	94.68	94.04	200.51	2.19


Table 2: Time (s) Taken to process various input file quantities

5 processes running in parallel


4.4x Speedup for training 250 images

2.2x Speedup for training 4000 images

Sequential vs. Parallelized



Sequential



Parallelized

Future Works

- Real time recognition
 - Optimize FaceX
- Hyperparameter tuning automation
- Integrated Preprocessing Code
 - Handle initialization calculation faster
 - Provide starter code in repo, tested to deliver identical DataPoints when compared to serial version




References

- [1] Cao X, Wei Y, Wen F, et al. Face alignment by explicit shape regression[J]. International Journal of Computer Vision, 2014, 107(2): 177-190.



Parallelizing Gradient Boosted Regression for Facial Landmark Recognition

Zhecheng Yao, Yixian Gan, Rebecca Qiu, Lucy Li
May 2023



Harvard John A. Paulson
School of Engineering
and Applied Sciences